

# classification-assignment-f23

October 26, 2023

0.1 Student name: Vangari Prashanth

0.2 Assignment: Regression

0.2.1 Part 1: Data Wrangling (50 pts)

You have to write code to answer the questions below 7 pts each subtask except for the first one (importing pandas...) which is worth 1 pt

```
[61]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

"""
Import pandas library
Read the data stored in your local machine https://www.kaggle.com/datasets/
    ↳ fedesoriano/stroke-prediction-dataset
Save data to a variable named df
"""
df = pd.read_csv('healthcare-dataset-stroke-data.csv') # Read the CSV file
df.head(5) # print the first five rows
```

```
[61]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	

	stroke
0	1
1	1


```
2      1
3      1
4      1
```

```
[62]: '''
      Use a pie chart to show the percentage of ever_married with labels and
      percentage rounded up to 2 decimals.
      '''

      count = df['ever_married'].value_counts()
      l = len(df)
      pcnt = (count / l) * 100
      pcnt = pcnt.round(2)

      label = ['Yes', 'No']

      plt.pie(pcnt, labels=label, autopct='%1.2f%%')
      plt.title('Percentage of Ever Married')
      plt.axis('equal')
      plt.show()
```



[63]:

```
'''
Encode the categorical columns to numeric. There are two types of encoding:
↳ ordinal and one-hot. Explain why you choose the encoding technique to the
↳ column(s) and implement it. Show some rows of df after encoding. There will
↳ be no printed console in this subtask
Reference (you may need incognito mode to browse the pages):
    https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.
↳ OrdinalEncoder.html
    https://towardsdatascience.com/
↳ guide-to-encoding-categorical-features-using-scikit-learn-for-machine-learning-5048997a5c79
    https://stackoverflow.com/questions/56502864/
↳ using-ordinalencoder-to-transform-categorical-values
    https://stackoverflow.com/questions/37292872/
↳ how-can-i-one-hot-encode-in-python
    https://pandas.pydata.org/docs/reference/api/pandas.get_dummies.html
'''
df_encoded = pd.get_dummies(df, columns=['gender', 'hypertension', 'age',
↳ 'heart_disease', 'avg_glucose_level', 'work_type', 'Residence_type',
↳ 'smoking_status'])
# df_encoded.head(5)

#Explanation:
#When there is no fundamental order or ranking in categorical data and each
↳ category is unique, one-hot encoding is used.
#When categories are ranked or have a meaningful order, ordinal encoding is
↳ used. It maintains the ordinal relationship by giving integer values to
↳ categories according to their respective locations.
#In this case, One Hot encoding is used because there is no natural order or
↳ ranking in the columns.
#For example, The gender column has male and female values, assigning 1 to male
↳ or female has no impact on the output.
```

[64]:

```
# Return boolean values indicating the number of missing rows of each column in
↳ ascending order. Do not overwrite df.

count = df.isna().sum()
srt = count.sort_values(ascending=True) # Sort the missing rows according to
↳ the values
p = (srt / len(df)) * 100 # Here, p is percentage
print(p)
```

id	0.000000
gender	0.000000
age	0.000000
hypertension	0.000000
heart_disease	0.000000
ever_married	0.000000

```
work_type          0.000000
Residence_type     0.000000
avg_glucose_level  0.000000
smoking_status     0.000000
stroke             0.000000
bmi                3.933464
dtype: float64
```

```
[65]: '''
Use one of imputation techniques in https://scikit-learn.org/stable/modules/impute.html#nearest-neighbors-imputation on bmi
Remember to keep df as of DataFrame type after applying imputation
Show the total number of missing values of the entire dataset
'''
from sklearn.impute import KNNImputer

# KNN Imputer creates a object with k=6
imputer = KNNImputer(n_neighbors=6)
df['bmi'] = imputer.fit_transform(df[['bmi']])
m = df.isna().sum().sum()
print("Total missing values in the dataset:", m)
```

Total missing values in the dataset: 0

```
[66]: '''Categorize bmi into groups as https://images.agoramedia.com/everydayhealth/gcms/BMI-in-Adults-722x406.jpg?width=722.
Print the column bmi after transformed. Do not overwrite to df
'''
# Categorize the weights based on their weights
counts = [(df['bmi'] < 18.5), (df['bmi'] >= 18.5) & (df['bmi'] <= 24.9), (df['bmi'] >= 25) & (df['bmi'] <= 29.9), (df['bmi'] >= 30) & (df['bmi'] <= 34.9), (df['bmi'] > 35)]
labels = ['Underweight', 'Normal', 'Overweight', 'Obese', 'Extremely Obese']

df_upd = df.assign(bmi_categories=np.select(counts, labels))

print(df_upd[['bmi_categories']])
```

```
      bmi_categories
0      Extremely Obese
1           Overweight
2              Obese
3              Obese
4             Normal
...              ...
5105          Overweight
5106  Extremely Obese
5107              Obese
```

```
5108      Overweight
5109      Overweight
```

```
[5110 rows x 1 columns]
```

```
[67]: '''
Show the column avg_glucose_level whose values are binned into 10 equal-sized
↳discrete intervals with labels such as "level 1", "level 2", etc.
Do not overwrite to df
'''
# avg_glucose_level binned to 10 equal sized intervals
df1 = df.assign(avg_glucose_level=pd.cut(df['avg_glucose_level'],bins=10,
    labels=['level 1', 'level 2', 'level 3', 'level 4', 'level 5', 'level_
↳6', 'level 7', 'level 8', 'level 9', 'level 10']))

print(df1['avg_glucose_level'])
```

```
0      level 9
1      level 7
2      level 3
3      level 6
4      level 6
...
5105    level 2
5106    level 4
5107    level 2
5108    level 6
5109    level 2
Name: avg_glucose_level, Length: 5110, dtype: category
Categories (10, object): ['level 1' < 'level 2' < 'level 3' < 'level 4' ...
'level 7' < 'level 8' <
                        'level 9' < 'level 10']
```

```
[68]: '''
Show the column avg_glucose_level whose values are binned into 5 equal_
↳percentile intervals as the output.
Do not overwrite to df
'''
df2 = df.assign(avg_glucose_level=pd.qcut(df['avg_glucose_level'],q=5))

print(df2['avg_glucose_level'])
```

```
0      (124.16, 271.74]
1      (124.16, 271.74]
2      (98.914, 124.16]
3      (124.16, 271.74]
4      (124.16, 271.74]
...
```

```

5105      (73.76, 85.6]
5106      (124.16, 271.74]
5107      (73.76, 85.6]
5108      (124.16, 271.74]
5109      (73.76, 85.6]
Name: avg_glucose_level, Length: 5110, dtype: category
Categories (5, interval[float64, right]): [(55.119, 73.76] < (73.76, 85.6] <
(85.6, 98.914] <
(98.914, 124.16] < (124.16, 271.74]]

```

## 0.2.2 Part 2: Logistic Regression(50 pts)

```

[69]: '''
10 pts:
Before implementing your LR model, you need to keep feature id since it is the
    ↳primary key to recognize patients.
However, your model cannot work (well) if it's kept as original of string type
    ↳while numeric type is meaningless.
One solution is to set it as index. In practice, you may have to have a
    ↳mechanism to convert it back to the original id for lookup.
Justify your solution.
'''

df = df.set_index('id')
df
print(df.dtypes)

from sklearn.impute import SimpleImputer
from sklearn.preprocessing import KBinsDiscretizer
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split

encoded_df = df
# Use numerical encoding to convert categorical data to numerical
encoded_df = pd.get_dummies(df, columns=['gender', 'ever_married', 'work_type',
    ↳'Residence_type', 'smoking_status'])

attrbts = ['age', 'hypertension', 'heart_disease', 'avg_glucose_level', 'bmi']
encoded_df[attrbts] = (encoded_df[attrbts] - encoded_df[attrbts].mean()) /
    ↳encoded_df[attrbts].std()

```

```

gender      object
age         float64
hypertension    int64
heart_disease  int64
ever_married  object
work_type     object

```

```

Residence_type      object
avg_glucose_level   float64
bmi                 float64
smoking_status      object
stroke              int64
dtype: object

```

```

[70]: '''
10 pts:
Assign X to be the whole df without column stroke and y to be the column stroke.
    ↳ Split X and y into X_train, X_test, y_train, and y_test with random_state=1,
    ↳ and test_size=0.2.
Should you use stratify? Explain
Reference: https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.train\_test\_split.html
    ↳ model_selection.train_test_split.html
'''

v = encoded_df['stroke'].value_counts()
print(v)

X = encoded_df.drop('stroke', axis=1)
y = encoded_df['stroke']

from sklearn.model_selection import train_test_split

# Split the data into training and test sets with stratification
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↳ random_state=1, stratify=y)

```

```

0    4861
1     249
Name: stroke, dtype: int64

```

20 pts Write a class My\_LR that implements Logistic Regression algorithm. You are required to have the following attributes \* Method: \* fit \* predict

Reference: [https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)  
**Using a pre-built library yields no credit. You have to write everything from scratch.**

```

[71]: class My_LR:

    def __init__(self, lr=0.01, rep=100000, intrcpt=True):
        self.lr = lr
        self.rep = rep
        self.intrcpt = intrcpt

    def __sgmd_func(self, z):
        return 1 / (1 + np.exp(-z))

```

```

def __intercept(self, X):
    i = np.ones((X.shape[0], 1))
    return np.concatenate((i, X), axis=1)

def fit(self, X, y):
    if self.intrcpt:
        X = self.__intercept(X)

    self.theta = np.zeros(X.shape[1])

    for i in range(self.rep):
        z = np.dot(X, self.theta)
        h = self.__sgmd_func(z)
        grdnt = np.dot(X.T, (h - y)) / y.size
        self.theta -= self.lr * grdnt

def predict(self, X):
    if self.intrcpt:
        X = self.__intercept(X)

    return np.round(self.__sgmd_func(np.dot(X, self.theta)))

```

```

[72]: # Run the code
reg = My_LR()
reg.fit(X_train, y_train)
y_pred = reg.predict(X_test)

print(y_pred)

```

[0. 0. 0. ... 0. 0. 0.]

```

[73]: '''
10 pts
Use a metric of either accuracy_score or balanced_accuracy_score to evaluate_
your predicted values and y_test
Explain why you prefer this metric over the other.
'''

from sklearn.metrics import accuracy_score, balanced_accuracy_score

# Accuracy score
acc = accuracy_score(y_test, y_pred)
print("Accuracy score:", acc)

# Balanced accuracy score
blncd_acc = balanced_accuracy_score(y_test, y_pred)

```



```
print("Balanced accuracy score:", blncd_acc)

'''
The accuracy score is helpful when classes are balanced, or when they account
↳for almost equal portions of the dataset.
Since it gives a clear view of the model's overall performance, the accuracy
↳score is a suitable metric to utilize in this case.
If the courses are not balanced, then the balanced accuracy score is a better
↳metric to use.
This is because it offers a more realistic picture of the model's output and
↳explains the underrepresentation of the minority class.
In general, it is important to consider the class balance and choose the metric
↳that is most appropriate for the given situation when determining a metric
↳to evaluate a model's performance.

'''
```

Accuracy score: 0.9510763209393346

Balanced accuracy score: 0.5

[73]: "\n\nThe accuracy score is helpful when classes are balanced, or when they account for almost equal portions of the dataset. \n\nSince it gives a clear view of the model's overall performance, the accuracy score is a suitable metric to utilize in this case. \n\nIf the courses are not balanced, then the balanced accuracy score is a better metric to use. \n\nThis is because it offers a more realistic picture of the model's output and explains the underrepresentation of the minority class. \n\nIn general, it is important to consider the class balance and choose the metric that is most appropriate for the given situation when determining a metric to evaluate a model's performance. \n\n\n"

[74]: #converting back the 'id' column from index to a column in the dataframe for  
↳lookup

```
df = df.reset_index()

print(df)
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	
...	...	...	...	...	...	...	
5105	18234	Female	80.0	1	0	Yes	
5106	44873	Female	81.0	0	0	Yes	
5107	19723	Female	35.0	0	0	Yes	

5108	37544	Male	51.0	0	0	Yes
5109	44679	Female	44.0	0	0	Yes

	work_type	Residence_type	avg_glucose_level	bmi	\
0	Private	Urban	228.69	36.600000	
1	Self-employed	Rural	202.21	28.893237	
2	Private	Rural	105.92	32.500000	
3	Private	Urban	171.23	34.400000	
4	Self-employed	Rural	174.12	24.000000	
...	...	...	...	...	
5105	Private	Urban	83.75	28.893237	
5106	Self-employed	Urban	125.20	40.000000	
5107	Self-employed	Rural	82.99	30.600000	
5108	Private	Rural	166.29	25.600000	
5109	Govt_job	Urban	85.28	26.200000	

	smoking_status	stroke
0	formerly smoked	1
1	never smoked	1
2	never smoked	1
3	smokes	1
4	never smoked	1
...	...	...
5105	never smoked	0
5106	never smoked	0
5107	never smoked	0
5108	formerly smoked	0
5109	Unknown	0

[5110 rows x 12 columns]